

Lab sheet 11 - Deep Learning

1. Neural Network Language Model

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 100)	500200
dense_2 (Dense)	(None, 100)	10100
dense_3 (Dense)	(None, 100)	10100
dense_4 (Dense)	(None, 5001)	505101

Total params: 1,025,501

Trainable params: 1,025,501

Non-trainable params: 0

Model fitting

```
81369/81369 [-----] - 104s 1ms/step - loss: 5.0930 - acc: 0.1934 - val_loss: 4.9334 - val_acc: 0.1646
1 - __newline__
Train on 81369 samples, validate on 23828 samples
Epoch 2/2
81369/81369 [-----] - 97s 1ms/step - loss: 5.0930 - acc: 0.1934 - val_loss: 4.9334 - val_acc: 0.1646
1 have , __newline__
Train on 81369 samples, validate on 23828 samples
Epoch 3/3
81369/81369 [-----] - 128s 2ms/step - loss: 4.8733 - acc: 0.2079 - val_loss: 4.8497 - val_acc: 0.1688
1 have , __newline__
Train on 81369 samples, validate on 23828 samples
Epoch 4/4
81369/81369 [-----] - 106s 1ms/step - loss: 4.7154 - acc: 0.2189 - val_loss: 4.8329 - val_acc: 0.1697
1 am a __unk__ , __newline__
Train on 81369 samples, validate on 23828 samples
Epoch 5/5
81369/81369 [-----] - 87s 1ms/step - loss: 4.5913 - acc: 0.2271 - val_loss: 4.8390 - val_acc: 0.1748
1 have , __newline__
Train on 81369 samples, validate on 23828 samples
Epoch 6/6
81369/81369 [-----] - 79s 976us/step - loss: 4.4088 - acc: 0.2322 - val_loss: 4.8048 - val_acc: 0.1725
1 have been absent friends , __newline__
Train on 81369 samples, validate on 23828 samples
Epoch 50/50
81369/81369 [-----] - 96s 1ms/step - loss: 3.6133 - acc: 0.2497 - val_loss: 6.5883 - val_acc: 0.1731
1 am a __unk__ , __newline__
```

Bigram model

```
81368/81368 [=====] - 83s 1ms/step - loss: 5.1627 - acc: 0.1859 - val_loss: 5.0539 - val_acc: 0.1548
i have not , __newline__
Train on 81368 samples, validate on 23827 samples
Epoch 3/3
81368/81368 [=====] - 59s 729us/step - loss: 4.8838 - acc: 0.2042 - val_loss: 4.9481 - val_acc: 0.1644
i have ' s s , __newline__
Train on 81368 samples, validate on 23827 samples
Epoch 4/4
81368/81368 [=====] - 61s 751us/step - loss: 4.6721 - acc: 0.2180 - val_loss: 4.9112 - val_acc: 0.1591
i have ' t . __newline__
Train on 81368 samples, validate on 23827 samples
Epoch 5/5
81368/81368 [=====] - 87s 1ms/step - loss: 4.5008 - acc: 0.2308 - val_loss: 4.8935 - val_acc: 0.1643
i have ' t . __newline__
Train on 81368 samples, validate on 23827 samples
Epoch 6/6
81368/81368 [=====] - 113s 1ms/step - loss: 4.3531 - acc: 0.2416 - val_loss: 4.9226 - val_acc: 0.1676
i have ' t . __newline__
Train on 81368 samples, validate on 23827 samples
Epoch 7/7
81368/81368 [=====] - 73s 899us/step - loss: 4.2191 - acc: 0.2485 - val_loss: 5.0076 - val_acc: 0.1656
i have ' t . __newline__
Train on 81368 samples, validate on 23827 samples
Epoch 8/8
81368/81368 [=====] - 73s 896us/step - loss: 4.0915 - acc: 0.2574 - val_loss: 5.0491 - val_acc: 0.1656

show more (open the raw output data in a text editor) ...

i have seen her die . __newline__
Train on 81368 samples, validate on 23827 samples
Epoch 50/50
81368/81368 [=====] - 103s 1ms/step - loss: 2.3393 - acc: 0.4445 - val_loss: 10.0073 - val_acc: 0.1415
i have not ended if this thou perceiv ' st , __newline__
```

Trigram model

```
Train on 81367 samples, validate on 23826 samples
Epoch 1/1
81367/81367 [=====] - 96s 1ms/step - loss: 5.7854 - acc: 0.1336 - val_loss: 5.2517 - val_acc: 0.1183
i have to , __newline__
Train on 81367 samples, validate on 23826 samples
Epoch 2/2
81367/81367 [=====] - 84s 1ms/step - loss: 5.2171 - acc: 0.1774 - val_loss: 5.1297 - val_acc: 0.1334
i have to not , __newline__
Train on 81367 samples, validate on 23826 samples
Epoch 3/3
81367/81367 [=====] - 76s 928us/step - loss: 4.9578 - acc: 0.1990 - val_loss: 5.0744 - val_acc: 0.1417
i have to have , __newline__
Train on 81367 samples, validate on 23826 samples
Epoch 4/4
81367/81367 [=====] - 78s 961us/step - loss: 4.7546 - acc: 0.2106 - val_loss: 5.0088 - val_acc: 0.1470
i have to not the __unk__ of __unk__ __newline__
Train on 81367 samples, validate on 23826 samples
Epoch 5/5
81367/81367 [=====] - 78s 955us/step - loss: 4.5719 - acc: 0.2243 - val_loss: 5.0878 - val_acc: 0.1560
i have to __unk__ , __newline__
Train on 81367 samples, validate on 23826 samples
Epoch 6/6
81367/81367 [=====] - 78s 858us/step - loss: 4.4026 - acc: 0.2368 - val_loss: 5.1442 - val_acc: 0.1579
i have to not the __unk__ of __unk__ , __newline__
Train on 81367 samples, validate on 23826 samples

show more (open the raw output data in a text editor) ...

i have to sworn me . __newline__
Train on 81367 samples, validate on 23826 samples
Epoch 50/50
81367/81367 [=====] - 48s 991us/step - loss: 1.6981 - acc: 0.5897 - val_loss: 12.6095 - val_acc: 0.1113
i have to beloved of what thou dost foist upon us that 's now manner women ' you th cold j ' s and he may be a tattered weed of small worth held : __newline__
```

2. Sentiment Classification using Neural Network

Model architecture

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 25)	375075
dropout (Dropout)	(None, 25)	0
dense_1 (Dense)	(None, 100)	2600
dropout_1 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 13)	663
softmax (Softmax)	(None, 13)	0

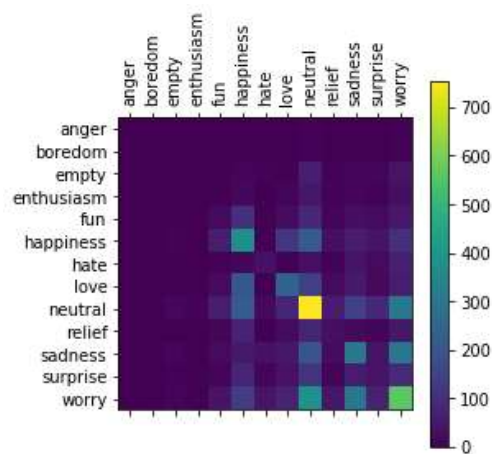
```
=====
```

Total params: 383,388
Trainable params: 383,388
Non-trainable params: 0

Test accuracy

```
250/250 [=====] - 7s 5ms/step - loss: 3.1381 - acc: 0.2965  
Test Set Accuracy: 29.65%
```

Confusion matrix



3. Per timestep prediction using Neural Network – POS tagging

Vocabulary building

```
Vocabulary size: 54562
Most frequent tokens
,: 173469
.: 138735
the: 109915
to: 68599
of: 64398
and: 59512
a: 57597
that: 48974
I: 44784
in: 40624
Least frequent tokens
incapacitates: 1
bankruptcies: 1
IPOs: 1
es: 1
Due: 1
Separate: 1
squashed: 1
raking: 1
Heroin: 1
```

Model Definition

```
Model: "sequential"

Layer (type)                 Output Shape              Param #
=====
embedding (Embedding)        (None, 100, 25)          250050
simple_rnn (SimpleRNN)         (None, 100, 30)          1680
time_distributed (TimeDistri (None, 100, 44)          1364
time_distributed_1 (TimeDist (None, 100, 44)          0
=====
Total params: 253,094
Trainable params: 253,094
Non-trainable params: 0
```

Test set accuracy

```
test_loss, test_acc = model.evaluate(X_test, y_test)
print(["Test Set Accuracy: %.2f%%"%(test_acc*100)])
✓ 22.1s Python
1312/1312 [=====] - 22s 11ms/step - loss: 0.0364 - acc:
0.9874
Test Set Accuracy: 98.74%
```

4. Pre trained embedding

Pre trained vectors

```
Loaded 7983/9999 pre-trained vectors
array([[ 0.          ,  0.          ,  0.          , ...,  0.          ,
         0.          ,  0.          ],
       [-0.010167 ,  0.020194 ,  0.21472999, ...,  0.18783   ,
        -0.84249997, -0.31208   ],
       [ 0.32543001, -0.089637 , -0.14733   , ..., -0.61750001,
        -0.54592001, -0.069893  ],
       ...,
       [-0.52920002,  0.90044999, -1.11520004, ...,  2.2046001 ,
        -0.87592   , -0.89977002],
       [-0.22378001, -0.088261 , -0.80461001, ...,  0.52086002,
        -0.024927 ,  0.35431999],
       [-0.93480998, -0.12461   , -0.44753   , ...,  0.32477   ,
        -0.53653997, -1.10179996]])
```

Training and testing accuracy

```
Train on 6722 samples, validate on 354 samples
Epoch 1/5
6722/6722 [=====] - 20s 3ms/step - loss: 2.2404 - acc: 0.4096 - val_loss: 2.1207 - val_acc: 0.4237
Epoch 2/5
6722/6722 [=====] - 15s 2ms/step - loss: 2.0493 - acc: 0.4528 - val_loss: 2.1413 - val_acc: 0.4181
Epoch 3/5
6722/6722 [=====] - 17s 3ms/step - loss: 2.0252 - acc: 0.4810 - val_loss: 2.3565 - val_acc: 0.3785
Epoch 4/5
6722/6722 [=====] - 18s 3ms/step - loss: 2.1230 - acc: 0.4311 - val_loss: 2.1379 - val_acc: 0.4294
Epoch 5/5
6722/6722 [=====] - 16s 2ms/step - loss: 2.0164 - acc: 0.4683 - val_loss: 2.1089 - val_acc: 0.4266
<keras.callbacks.callbacks.History at 0x21fabac3ec8>
```

```
loss, acc = model.evaluate(X_test, to_categorical(y_test))
print("Test accuracy: %0.2f%%"%(acc*100))
```

```
1770/1770 [=====] - 2s 877us/step
Test accuracy: 48.87%
```

5. Recurrent Neural Network

Data Loading

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/reuters.npz
2113536/2110848 [=====] - 1s 0us/step
C:\Users\Kavishka\anaconda3\lib\site-
```

Data preparing

```
[[ 0  0  0 ... 15 17 12]
 [ 0  0  0 ... 505 17 12]
 [ 0  0  0 ... 11 17 12]
 ...
 [ 0  0  0 ... 254 17 12]
 [ 0  0  0 ... 2735 17 12]
 [ 0  0  0 ... 4329 17 12]]
```

Simple RNN

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 200, 10)	100000
simple_rnn_1 (SimpleRNN)	(None, 25)	900
dense_1 (Dense)	(None, 46)	1196
softmax_1 (Softmax)	(None, 46)	0

Total params: 102,096

Trainable params: 102,096

Non-trainable params: 0

Result

Epoch 10/10

```
132/132 [=====] - 6s 42ms/step - loss: 1.1863 - acc: 0.6944 -
val_loss: 2.0224 - val_acc: 0.5270
```

```
139/139 [=====] - 2s 15ms/step - loss: 2.2730 - acc: 0.4872
```

Test accuracy: 48.72%

Bidirectional RNN

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, 200, 10)	100000

bidirectional (Bidirectional)	(None, 25)	1800

dense_2 (Dense)	(None, 46)	1196

softmax_2 (Softmax)	(None, 46)	0
=====		

Total params: 102,996

Trainable params: 102,996

Non-trainable params: 0

Result

- val_loss: 1.9244 - val_acc: 0.4820

139/139 [=====] - 3s 19ms/step - loss: 2.2139 - acc: 0.4140

Test accuracy: 41.40%

Simple RNN with averaging

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
embedding_3 (Embedding)	(None, 200, 10)	100000

simple_rnn_3 (SimpleRNN)	(None, 200, 25)	900

global_average_pooling1d (GlobalAveragePooling1D)	(None, 25)	0

dense_3 (Dense)	(None, 46)	1196

softmax_3 (Softmax)	(None, 46)	0
=====		

Total params: 102,096

Trainable params: 102,096

Non-trainable params: 0

Result

```
Epoch 5/5
132/132 [=====] - 9s 71ms/step - loss: 1.9451 - acc: 0.4823 -
val_loss: 1.5894 - val_acc: 0.5901
139/139 [=====] - 3s 22ms/step - loss: 1.8558 - acc: 0.5168
Test accuracy: 51.68%
```

Multilayer Bidirectional RNN

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
embedding_4 (Embedding)	(None, 200, 10)	100000

bidirectional_1 (Bidirection	(None, 200, 25)	7200

bidirectional_2 (Bidirection	(None, 25)	10200

dense_4 (Dense)	(None, 46)	1196

softmax_4 (Softmax)	(None, 46)	0
=====		
Total params: 118,596		
Trainable params: 118,596		
Non-trainable params: 0		

Result

```
Epoch 5/5
132/132 [=====] - 23s 172ms/step - loss: 2.2263 - acc: 0.4102
- val_loss: 1.9258 - val_acc: 0.4820
139/139 [=====] - 6s 44ms/step - loss: 2.2125 - acc: 0.4140
Test accuracy: 41.40%
```


6. Reuters News Classification

Data loading

```
Train on 6722 samples, validate on 354 samples
Epoch 1/5
6722/6722 [=====] - 20s 3ms/step - loss: 2.2404 - acc: 0.4096 - val_loss: 2.1207 - val_acc: 0.4237
Epoch 2/5
6722/6722 [=====] - 15s 2ms/step - loss: 2.0493 - acc: 0.4528 - val_loss: 2.1413 - val_acc: 0.4181
Epoch 3/5
6722/6722 [=====] - 17s 3ms/step - loss: 2.0252 - acc: 0.4810 - val_loss: 2.3565 - val_acc: 0.3785
Epoch 4/5
6722/6722 [=====] - 18s 3ms/step - loss: 2.1230 - acc: 0.4311 - val_loss: 2.1379 - val_acc: 0.4294
Epoch 5/5
6722/6722 [=====] - 16s 2ms/step - loss: 2.0164 - acc: 0.4683 - val_loss: 2.1089 - val_acc: 0.4266

<keras.callbacks.callbacks.History at 0x21fabac3ec8>
```

```
loss, acc = model.evaluate(X_test, to_categorical(y_test))
print("Test accuracy: %0.2f%%"%(acc*100))
```

```
1770/1770 [=====] - 2s 877us/step
Test accuracy: 48.87%
```

Dense NN model

```
Train on 6722 samples, validate on 354 samples
Epoch 1/5
6722/6722 [=====] - 20s 3ms/step - loss: 2.2404 - acc: 0.4096 - val_loss: 2.1207 - val_acc: 0.4237
Epoch 2/5
6722/6722 [=====] - 15s 2ms/step - loss: 2.0493 - acc: 0.4528 - val_loss: 2.1413 - val_acc: 0.4181
Epoch 3/5
6722/6722 [=====] - 17s 3ms/step - loss: 2.0252 - acc: 0.4810 - val_loss: 2.3565 - val_acc: 0.3785
Epoch 4/5
6722/6722 [=====] - 18s 3ms/step - loss: 2.1230 - acc: 0.4311 - val_loss: 2.1379 - val_acc: 0.4294
Epoch 5/5
6722/6722 [=====] - 16s 2ms/step - loss: 2.0164 - acc: 0.4683 - val_loss: 2.1089 - val_acc: 0.4266

<keras.callbacks.callbacks.History at 0x21fabac3ec8>
```

```
loss, acc = model.evaluate(X_test, to_categorical(y_test))
print("Test accuracy: %0.2f%%"%(acc*100))
```

```
1770/1770 [=====] - 2s 877us/step
Test accuracy: 48.87%
```

Vanilla RNN

```
Train on 5572 samples, validate on 620 samples
Epoch 1/10
5572/5572 [=====] - 18s 3ms/step - loss: 2.7791 - acc: 0.2859 - val_loss: 2.2163 - val_acc: 0.4145
Epoch 2/10
5572/5572 [=====] - 18s 3ms/step - loss: 2.2394 - acc: 0.4072 - val_loss: 2.1670 - val_acc: 0.4145
Epoch 3/10
5572/5572 [=====] - 19s 3ms/step - loss: 2.8754 - acc: 0.3173 - val_loss: 2.2581 - val_acc: 0.4145
Epoch 4/10
5572/5572 [=====] - 15s 3ms/step - loss: 2.4480 - acc: 0.3568 - val_loss: 2.2328 - val_acc: 0.4145
Epoch 5/10
5572/5572 [=====] - 14s 3ms/step - loss: 2.5381 - acc: 0.3487 - val_loss: 3.2278 - val_acc: 0.3871
Epoch 6/10
5572/5572 [=====] - 15s 3ms/step - loss: 2.3346 - acc: 0.3771 - val_loss: 2.2153 - val_acc: 0.4145
Epoch 7/10
5572/5572 [=====] - 15s 3ms/step - loss: 2.7480 - acc: 0.3320 - val_loss: 3.2958 - val_acc: 0.1903
Epoch 8/10
5572/5572 [=====] - 15s 3ms/step - loss: 2.7328 - acc: 0.3186 - val_loss: 2.2404 - val_acc: 0.4145
Epoch 9/10
5572/5572 [=====] - 17s 3ms/step - loss: 2.2921 - acc: 0.4013 - val_loss: 2.2340 - val_acc: 0.4145
Epoch 10/10
5572/5572 [=====] - 17s 3ms/step - loss: 2.2890 - acc: 0.3993 - val_loss: 2.2298 - val_acc: 0.4145
2654/2654 [=====] - 2s 665us/step
Test accuracy: 41.07%
```

Bidirectional RNN

```
Train on 5572 samples, validate on 620 samples
Epoch 1/10
5572/5572 [=====] - 15s 3ms/step - loss: 2.4509 - acc: 0.3835 - val_loss: 2.2217 - val_acc: 0.4145
Epoch 2/10
5572/5572 [=====] - 15s 3ms/step - loss: 2.2142 - acc: 0.4139 - val_loss: 2.1106 - val_acc: 0.4226
Epoch 3/10
5572/5572 [=====] - 19s 3ms/step - loss: 2.0135 - acc: 0.4883 - val_loss: 1.8701 - val_acc: 0.5403
Epoch 4/10
5572/5572 [=====] - 19s 3ms/step - loss: 1.8171 - acc: 0.5558 - val_loss: 1.8466 - val_acc: 0.5435
Epoch 5/10
5572/5572 [=====] - 19s 3ms/step - loss: 1.6924 - acc: 0.5811 - val_loss: 1.8101 - val_acc: 0.5597
Epoch 6/10
5572/5572 [=====] - 18s 3ms/step - loss: 1.6138 - acc: 0.5951 - val_loss: 1.8306 - val_acc: 0.5210
Epoch 7/10
5572/5572 [=====] - 20s 4ms/step - loss: 1.5124 - acc: 0.6129 - val_loss: 1.7926 - val_acc: 0.5565
Epoch 8/10
5572/5572 [=====] - 19s 3ms/step - loss: 1.4526 - acc: 0.6346 - val_loss: 1.7773 - val_acc: 0.5855
Epoch 9/10
5572/5572 [=====] - 20s 4ms/step - loss: 1.3503 - acc: 0.6630 - val_loss: 1.7788 - val_acc: 0.5645
Epoch 10/10
5572/5572 [=====] - 19s 3ms/step - loss: 1.2750 - acc: 0.6843 - val_loss: 1.8269 - val_acc: 0.5500
2654/2654 [=====] - 2s 798us/step
Test accuracy: 57.87%
```

Multilayer Bidirectional RNN

```

Train on 5572 samples, validate on 620 samples
Epoch 1/5
5572/5572 [=====] - 55s 10ms/step - loss: 2.6555 - acc: 0.3589 - val_loss: 2.2196 - val_acc: 0.4145
Epoch 2/5
5572/5572 [=====] - 68s 12ms/step - loss: 2.3062 - acc: 0.4034 - val_loss: 2.2078 - val_acc: 0.4145
Epoch 3/5
5572/5572 [=====] - 65s 12ms/step - loss: 2.1294 - acc: 0.4212 - val_loss: 1.8419 - val_acc: 0.4629
Epoch 4/5
5572/5572 [=====] - 97s 17ms/step - loss: 1.7636 - acc: 0.5422 - val_loss: 1.6170 - val_acc: 0.5758
Epoch 5/5
5572/5572 [=====] - 67s 12ms/step - loss: 1.6563 - acc: 0.5659 - val_loss: 1.5789 - val_acc: 0.5532
2654/2654 [=====] - 7s 3ms/step
Test accuracy: 54.97%

```

7. Text-Generation-with-LSTM

```
Epoch 3/3
20020/200201 [-----] - 25% 1m/step - loss: 1.9703
--- Generating with seed: "on that dream as reality, that is, we
accept our hypotheses "
----- temperature: 0.3
on that dream as reality, that is, we
accept our hypotheses and the string the simplifying the problem of the consciousness of the words of the prescription of the presence of the fact the consciousness of the believes the more the prescription of
----- temperature: 0.5
a the string to the preserving and the simplifying of the simplifying and are that the really simplifying of the begrow the consciousness of the person of the master of the consciousness of the fore say the also to
----- temperature: 1.0
conjunction of the happiness of the more of the richness of richness of still wort ever-the alienation.

|   an ideal
vowshere, lovelies to odious,
to be vitality and for elax
him is grifonious odious--the simplifying him every
higher way of antigals
find that their twined of avoid of the ruling
insure by better viable strindled--of allow to
sublestradich well and doors, tropical feacotige disconformance an in eat theire of liense green trocivigal , t
----- temperature: 1.2
disconformance an in eat theire of liense green trocivigal , the y.

Ehestinists with militities, we sons of gunn' idemosty,
nothing odious--stratid gatarality
wastely outrage--no best perhaps, the nodes article
senury
also-appending life of the spirit, portion to be after
woties greety is seedif
```

```

titles
is
literally therey frok; you spurn from t
----- temperature: 1.1
y is howifilities
is
literally therey frok; you spurn from the scientific most comprehensive the sense degree of these, in
encouraging, the parasite
and
philosophy "these you discernment"---knowledge, however: it without of so-racily should, is actual
stuff-indagation,
withawareness of historicalism by the race of a instincts; and there est to
i back), (guiltlessen any serious grows that the ty of his beyoon
1 : e; well
according a subercompai
epoch 11
epoch 1/1
200201/200201 [-----] - 100k bps/step - loss: 1.000
--- Generating with seed: "gent in reaching man; and
before it has arrived man denies--"
----- temperature: 0.1
gent in reaching man; and
before it has arrived man denies--the the enemy and such a man who has a still in the suspect of the suspect of the spirit and such a reports itself is the sense of the sense of the sense of the sense
time to the sense of the present and anders of the sense of the sense of the present upon the same the sense of the same and desire to the person to the sense of the sense of the sense of the s
----- temperature: 0.1
----- temperature: 0.1
----- temperature: 0.1
we turn
epoch 11
progress 11
"particula nothing more schoolen but clittering_ aswell with lives jangliterations and will
act rather have means of
low indiv

```